

1. Discuss Public, Private, and Hybrid Clouds

Cloud computing allows sharing of computing resources using the Internet. Based on how resources are deployed and accessed, cloud models are categorized into **Public**, **Private**, and **Hybrid Clouds**.

1. Public Cloud

- A **public cloud** is owned and operated by third-party service providers like **Amazon AWS, Google Cloud, or Microsoft Azure**.
 - It is built **over the Internet** and is **accessible to any user** who pays for the service.
 - The provider manages everything – infrastructure, storage, and platforms – and the user only needs to **subscribe and access**.
 - It offers a **pay-per-use pricing model**, which makes it **cost-effective for startups and individuals**.
 - Public clouds are **geographically distributed** to reduce response times and improve fault tolerance.
 - Examples: **Google App Engine, AWS EC2, Microsoft Azure, IBM Blue Cloud**.
-

2. Private Cloud

- A **private cloud** is built and managed **within an organization's intranet**.
 - It is **owned and operated by a single organization**, and access is restricted to that organization and its trusted partners.
 - It is not made for public use or sale and focuses on **security, control, and customization**.
 - Private clouds offer **higher efficiency, privacy, and reliability** but come with **higher setup and maintenance costs**.
 - They are best for companies that deal with **sensitive data**, such as banks, healthcare, and government agencies.
-

3. Hybrid Cloud

- A **hybrid cloud** combines **both public and private clouds** to leverage the benefits of both.
- Organizations can keep **sensitive data in the private cloud** and use the **public cloud for less critical workloads**.
- It enables **flexibility and scalability**. For example, if the private cloud runs out of resources, it can **burst into the public cloud**.

- Hybrid clouds are suitable for businesses with **dynamic workloads** and **compliance requirements**.
- Example: **IBM Research Compute Cloud (RC2)** interconnects multiple research centers using a hybrid model.

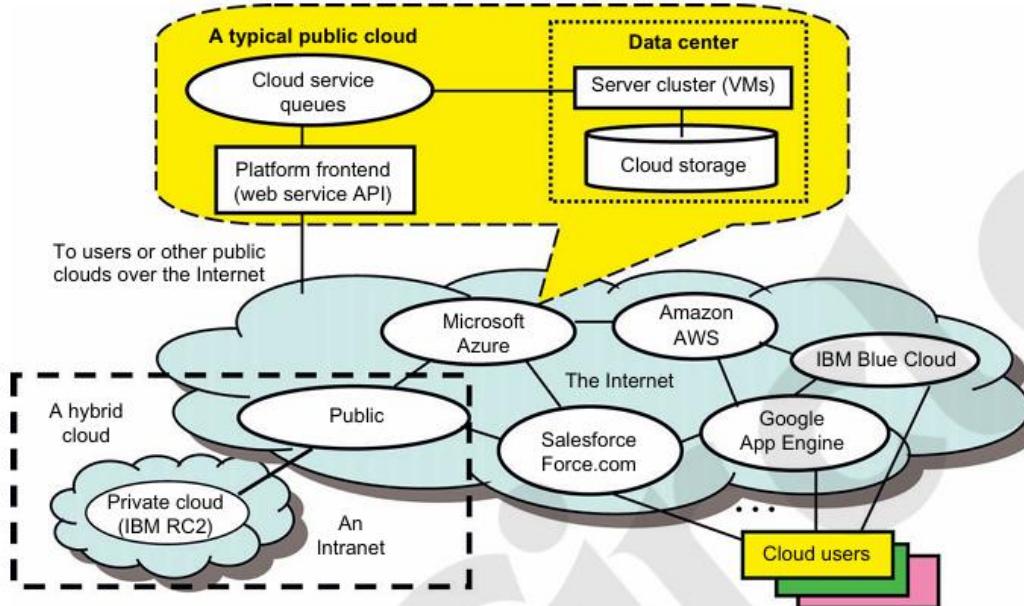
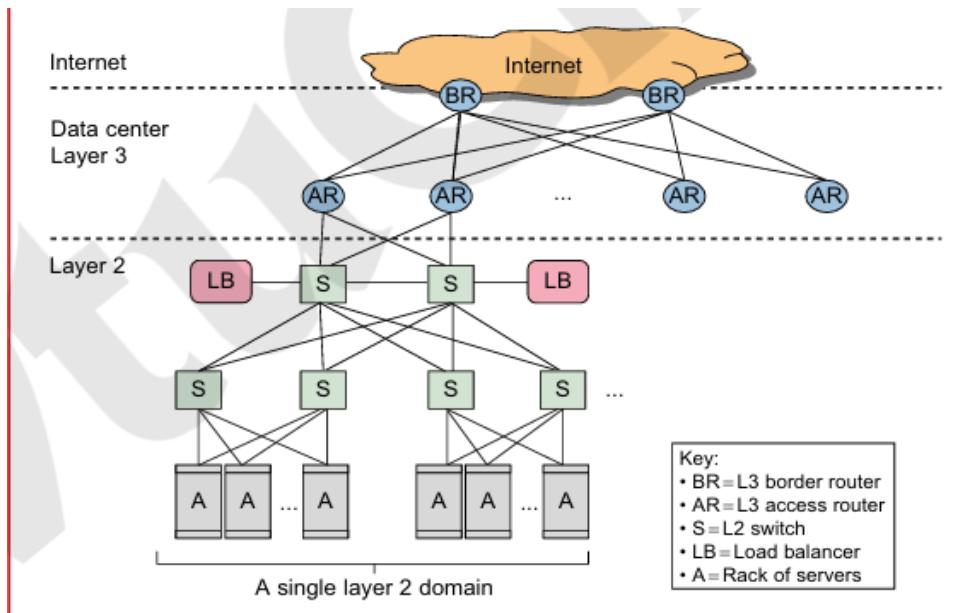


FIGURE 4.1

Public, private, and hybrid clouds illustrated by functional architecture and connectivity of representative clouds available by 2011.

2. Explain Data-Center Networking Structure



🌐 Overview of Data-Center Networking in Cloud Computing

- The **core** of a cloud system is made of **server clusters** (or VM clusters).

- **Compute nodes** in the cluster handle user jobs, while **control nodes** manage and monitor system activities.
 - **Gateway nodes** serve as the entry point for users and also manage **security** and access control.
-

Role of Virtual Clusters and Job Scheduling

- User jobs are assigned to **virtual clusters**, which are created on-demand from physical resources.
- These clusters are **dynamic**, meaning they grow or shrink based on workload requirements.
- This helps the cloud system handle **fluctuating workloads** efficiently.

Multilayer Network Architecture

- **Layer 2 (Bottom Layer)**: Contains **server racks** connected by **fast switches (S)**.
- **Layer 3 (Top Layer)**: Contains **access routers (AR)** and **border routers (BR)** that link the data center to the **Internet**.
- This layered design ensures **high-speed internal communication** and **smooth external access**.

Difference Between Traditional and Cloud-Based Clusters

- In traditional clusters or grids, resource demand is usually **static** and predictable.
 - Cloud systems, especially **private clouds**, are designed for **dynamic resource allocation** to meet changing user demands.
-

Scale and Structure of Data Centers

- Data centers are **highly scalable**, often consisting of **thousands to millions** of servers.
- For example, Microsoft's data center in Chicago has **100,000 eight-core servers** across **50 containers**.
- Unlike supercomputers, data centers use **disk-based storage** on server nodes, **memory cache**, and **databases** for storage.

3. Explain Cloud Ecosystems

Cloud Ecosystem – The Big Picture The **cloud ecosystem** includes **providers**, **users**, and a variety of **cloud-enabling tools** and **technologies**.

- This ecosystem emerged strongly with **public clouds**, but also supports **private** and **hybrid** clouds.

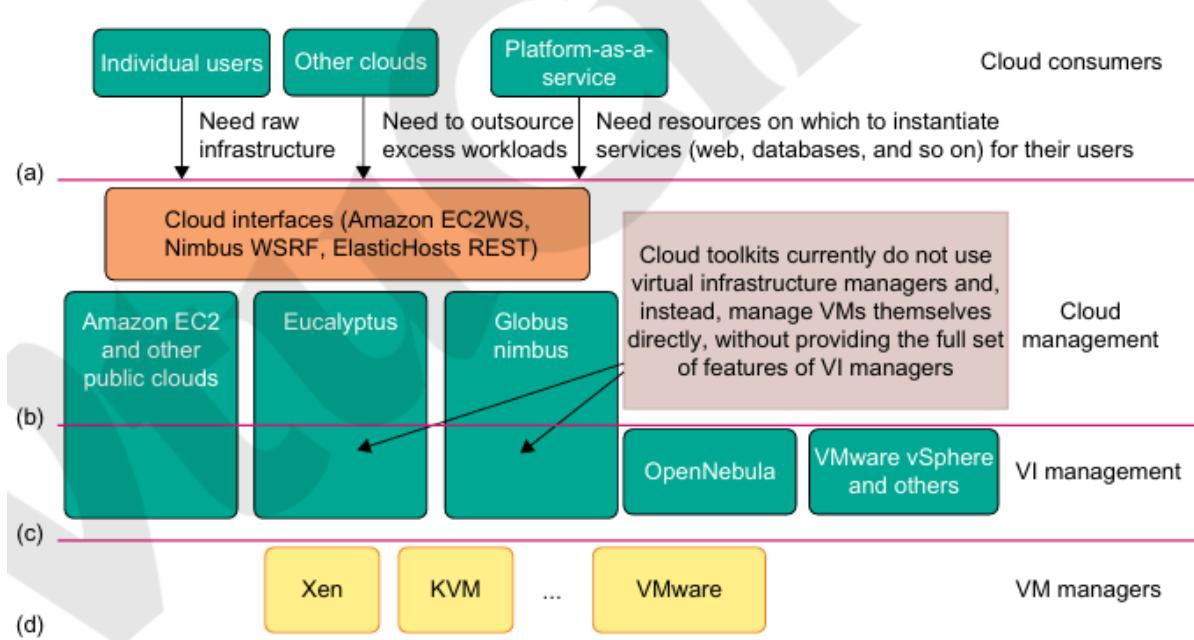


FIGURE 4.4

Cloud ecosystem for building private clouds: (a) Consumers demand a flexible platform; (b) Cloud manager

💡 2. Level 1: Cloud Users / Consumers (Top Layer)

- At the **top of the ecosystem**, **end-users or enterprises** interact with cloud services.
- These users expect a **flexible**, scalable, and on-demand platform — typically via **web service interfaces** like:
 - Amazon EC2 API
 - Nimbus WSRF
 - ElasticHost REST APIs
- They may access both **public cloud resources** or **remote services** exposed by **private/hybrid clouds**.

🛠 3. Level 2: Cloud Management Layer

- This is the **IaaS-level layer** where **cloud managers** virtualize and expose physical resources to users.
- It provides **on-demand provisioning** of compute/storage/networking via:
 - Amazon EC2
 - Eucalyptus (open source)
 - Nimbus (open source)
- The goal here is **resource abstraction**, **self-service**, and **elasticity**.

4. Level 3: Virtual Infrastructure (VI) Management Layer

- VI management tools **allocate VMs** across **clusters of physical servers**.
 - They manage **infrastructure automation**, such as:
 - **Server consolidation**
 - **Load balancing**
 - **Infrastructure partitioning and resizing**
 - Tools used in this layer include:
 - **OpenNebula**
 - **VMware vSphere**
 - **oVirt**
 - **Platform VM Orchestrator**
 - These tools act as a **bridge** between cloud-level provisioning and physical infrastructure.
-

5. Level 4: VM Management Layer (Bottom Layer)

- At the lowest level, **VM managers** control the **deployment, execution, and lifecycle** of VMs on **individual host machines**.
- These host-level operations include:
 - **Starting/stopping VMs**
 - **Monitoring VM health**
 - **Allocating host-level resources**
- Underlying hypervisors like **Xen**, **KVM**, and **VMware** are controlled via VI tools.

4. Discuss the case scenario of Google App Engine (GAE) for Platform as a Service (PaaS) application.

1. What is GAE?

- **Google App Engine (GAE)** is a classic example of the **Platform as a Service (PaaS)** model.
- It allows developers to **build, test, deploy, and manage web applications** on Google's cloud infrastructure.
- GAE abstracts away the hardware, OS, and networking, so developers can **focus only on app logic and functionality**.

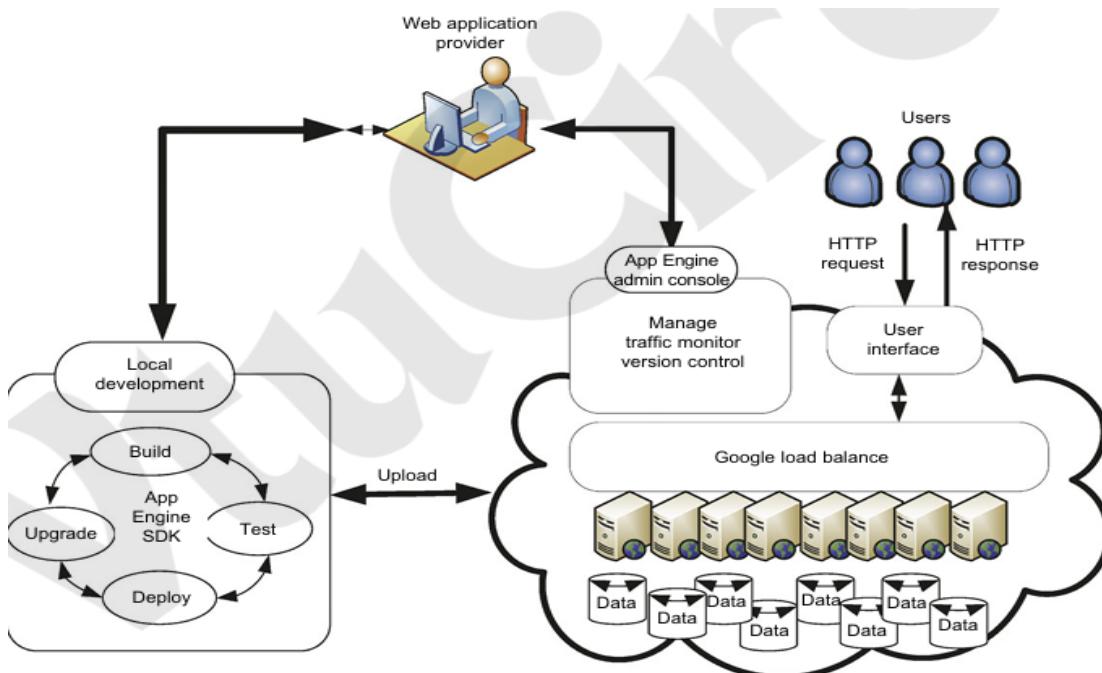


FIGURE 4.7

Google App Engine platform for PaaS operations.

◆ 2. How GAE Works – The Operational Model

- Applications run on **Google's powerful server clusters**, and users **share infrastructure** with other apps.
- GAE handles:
 - **Automatic scaling** (based on traffic)
 - **Load balancing** (even distribution of requests)
 - **Task scheduling** (events triggered at specific times or intervals)

◆ 3. Development Flow in GAE

1. Local Development:

- Google provides a **fully featured SDK** that simulates the entire GAE environment on the developer's local machine.
- Developers can **code, test, and debug locally**, just like in traditional software development.

2. Deployment:

- After development and testing, the application is **uploaded directly to Google's infrastructure** using GAE deployment tools.
- The cloud handles all hosting, resource allocation, scaling, and updates automatically.

◆ 4. Supported Tools & Languages

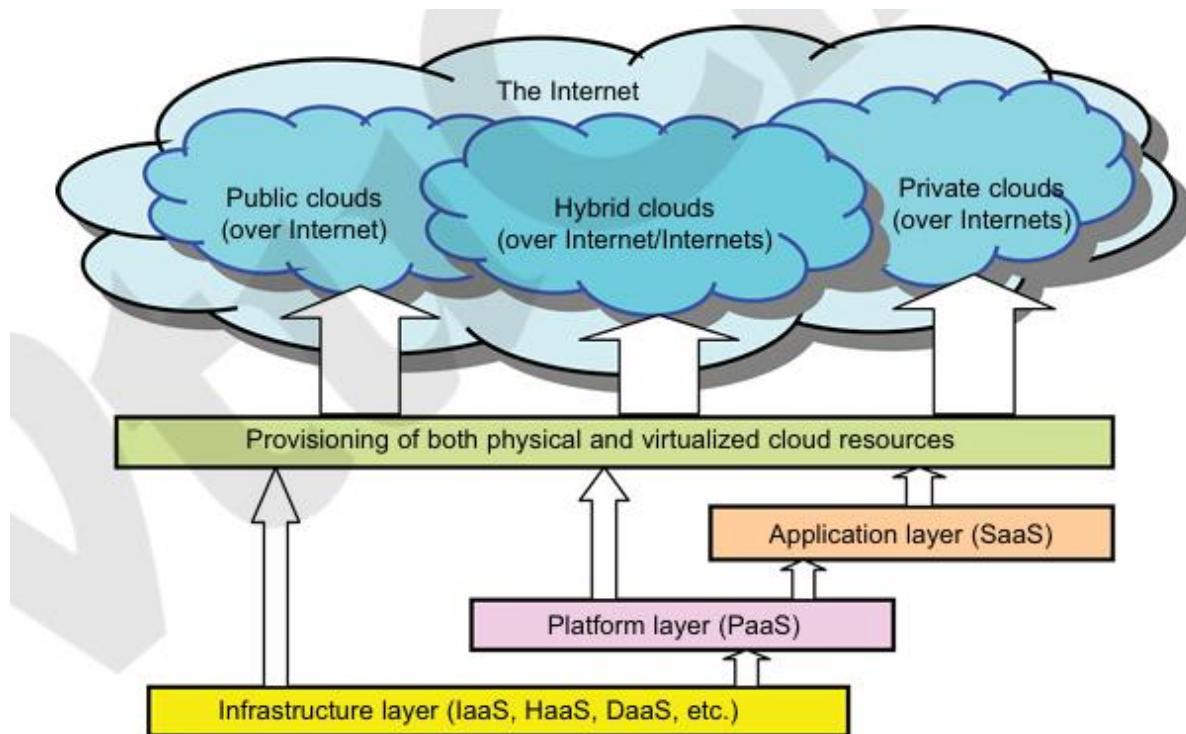
- GAE supports various programming languages and tools, such as **Python, Java, .NET**, etc.
 - It provides **APIs, runtime libraries, and software development environments** suited for building modern web apps.
-

◆ 5. PaaS Benefits via GAE

- **No infrastructure management** required.
- Facilitates **collaborative development** by distributed teams.
- Supports **third-party tools** for software management, integration, and monitoring.
- Ideal for **fast prototyping, scalable web apps, and startups**.

5. Discuss Layered Cloud Architectural Development for IaaS, PaaS, and SaaS

Cloud architectures are structured in **three main layers: Infrastructure, Platform, and Application**. These layers provide a clear division of responsibilities between the cloud provider and the user, and the services are built on top of each other to provide a comprehensive cloud experience. Here's an explanation of how each layer works:



1. Infrastructure Layer (IaaS)

- **Key Concept:** The **Infrastructure Layer** provides the fundamental hardware and networking resources that form the foundation of any cloud service.
 - **Components:**
 - **Compute resources:** Virtual machines (VMs), containers.
 - **Storage resources:** Block storage, object storage, databases.
 - **Networking:** Virtual networks, VPNs, load balancers, DNS.
 - **Role in Cloud Architecture:**
 - It is **the starting point** for cloud services (IaaS).
 - In **IaaS**, the cloud provider offers **virtualized computing resources** and **networking resources** (compute, storage, and networking).
 - Users can **rent resources** and manage the underlying infrastructure, which includes running virtual machines, setting up networks, and configuring storage.
 - **Virtualization** and **automation** are key in this layer to dynamically provision resources, optimize usage, and scale based on demand.
 - **Example: Amazon EC2** — Provides virtualized CPU and storage resources, but users manage their own operating systems, middleware, and applications.
-

2. Platform Layer (PaaS)

- **Key Concept:** The **Platform Layer** abstracts much of the complexity of infrastructure management and provides tools and frameworks to develop and deploy applications.
- **Components:**
 - **Development environments:** Integrated development tools, programming frameworks.
 - **Databases and middleware:** Managed databases (e.g., SQL, NoSQL), messaging queues, caching services.
 - **Application monitoring:** Tools to test, monitor, and analyze applications.
 - **Security and scalability:** Automated scaling, security patches, performance monitoring.
- **Role in Cloud Architecture:**
 - The **PaaS layer** allows developers to focus on writing code and deploying applications without worrying about managing the underlying infrastructure (virtual machines, networking, storage, etc.).

- The platform provides **middleware**, **software libraries**, and pre-configured environments, which make it easier to develop, deploy, and maintain applications.
 - It includes **automation** for scaling applications, **load balancing**, **version control**, and integrated services for testing and monitoring application performance.
 - **Example: Google App Engine** — Provides a runtime environment for developing web apps in languages like Java, Python, and PHP, where developers only manage their app code and the platform handles everything else (scaling, provisioning, load balancing).
-

3. Application Layer (SaaS)

- **Key Concept:** The **Application Layer** is where cloud services are delivered to users as fully managed software applications.
 - **Components:**
 - **User-facing applications:** These are the end-user services that are ready to be used out of the box (e.g., email, CRM, office suites).
 - **Enterprise software solutions:** Business management tools like **CRM**, **ERP**, **financial services**, **e-commerce platforms**.
 - **Integration tools:** APIs and other tools to integrate SaaS apps with other platforms and data sources.
 - **Role in Cloud Architecture:**
 - The **SaaS layer** delivers **complete, fully-managed software solutions** to the end user.
 - Users do **not need to manage infrastructure** or platforms; they simply use the software directly.
 - The provider manages **everything**: infrastructure, platform, updates, scalability, security.
 - **SaaS services** are typically multi-tenant and provide high availability, security, and auto-scaling.
 - **Example: Salesforce.com** — A CRM software where the provider manages everything, from infrastructure to platform, and users only interact with the application to manage customer relationships, marketing, and sales.
-

4. Interrelationship of Layers

- **Layered Dependency:** Each layer in cloud architecture builds upon the one below it.

- **IaaS** serves as the foundation, providing the virtualized infrastructure for PaaS.
- **PaaS** leverages IaaS to create a platform where developers can create applications without worrying about the underlying hardware.
- **SaaS** sits at the top, offering fully functional applications to end users.
- **Layer Interaction:**
 - **IaaS** is primarily for **resource provisioning** and gives users more control over the infrastructure.
 - **PaaS** offers an **environment for development** with tools, frameworks, and services for easier app deployment and management.
 - **SaaS** delivers **complete applications** that users can access without any technical management.

6. Extended Cloud Computing Services

In cloud computing, the **extended services** go beyond the traditional layers of SaaS, PaaS, and IaaS. These extended services provide additional functionalities that enhance cloud offerings by supporting physical infrastructure, network interconnections, data, security.

Cloud application (SaaS)			Concur, RightNOW, Teleo, Kenexa, Webex, Blackbaud, salesforce.com, Netsuite, Kenexa, etc.
Cloud software environment (PaaS)			Force.com, App Engine, Facebook, MS Azure, NetSuite, IBM BlueCloud, SGI Cyclone, eBay
Cloud software infrastructure			Amazon AWS, OpSource Cloud, IBM Ensembles, Rackspace cloud, Windows Azure, HP, Banknorth
Computational resources (IaaS)	Storage (DaaS)	Communications (Caas)	Savvis, Internap, NTTCommunications, Digital Realty Trust, 365 Main
Collocation cloud services (LaaS)			Owest, AT&T, AboveNet
Network cloud services (NaaS)			VMware, Intel, IBM, XenEnterprise
Hardware/Virtualization cloud services (HaaS)			

1. Hardware as a Service (HaaS)

- **Key Concept:** This is the **bottom-most layer** that provides physical hardware resources such as servers, storage, and other physical devices as a service.
- **Components:**
 - **Computing hardware:** Physical servers, storage devices, and networking equipment.
 - **Dedicated hardware:** Servers that are dedicated to a particular customer or workload.

- **Example:** Data centers offering the physical machines used to run virtualized environments on IaaS platforms, like **Amazon EC2** or **Microsoft Azure Virtual Machines**.
-

2. Network as a Service (NaaS)

- **Key Concept:** NaaS refers to the **virtualized network infrastructure** provided as a service. It connects hardware components in the cloud and provides networking features such as routing, firewalls, load balancing, and VPNs.
 - **Components:**
 - **Virtual Networks (VLANs):** Virtual LANs that connect cloud resources securely.
 - **Firewall as a Service:** Virtualized firewall solutions.
 - **Virtual Private Network (VPN):** Secure, private connections between different cloud environments or between cloud and on-premises systems.
 - **Example:** **Amazon Virtual Private Cloud (VPC)** allows users to create isolated networks within the AWS cloud, along with routing, subnetting, and firewall management.
-

3. Location as a Service (LaaS)

- **Key Concept:** LaaS refers to services related to **collocation**, providing physical data centers or spaces for users to host their computing infrastructure.
 - **Components:**
 - **Collocation services:** Hosting services that allow organizations to place their hardware in secure data centers.
 - **Power, cooling, and security:** Physical infrastructure that ensures data center uptime and protection.
 - **Example:** **Equinix** provides data center colocation services where clients can place their hardware and manage it securely.
-

4. Data as a Service (DaaS)

- **Key Concept:** DaaS refers to **on-demand data storage and processing services** in the cloud.
- **Components:**
 - **Cloud-based storage:** Offers scalable storage options for data, such as object storage (e.g., Amazon S3).

- **Data integration services:** Helps users connect, synchronize, and transfer data between cloud services and third-party applications.
 - **Example:** **Google Cloud Storage** offers scalable, secure data storage that can be used across various cloud applications.
-

5. Communication as a Service (CaaS)

- **Key Concept:** CaaS provides cloud-based communication tools for voice, video, and messaging services.
 - **Components:**
 - **Voice over IP (VoIP):** Internet-based phone systems.
 - **Video conferencing:** Virtual communication tools.
 - **Messaging platforms:** Cloud-based messaging systems for team collaboration.
 - **Example:** **Twilio** provides a platform for integrating SMS, voice, and video communication capabilities into applications.
-

6. Security as a Service (SecaaS)

- **Key Concept:** SecaaS refers to cloud-based **security services** that help organizations protect their data and applications.
- **Components:**
 - **Identity and Access Management (IAM):** Services for controlling user access and authentication.
 - **Intrusion Detection and Prevention Systems (IDPS):** Tools for detecting and preventing attacks.
 - **Encryption:** Protecting data both in transit and at rest.
- **Example:** **Cloudflare** provides cloud security services such as DDoS protection, web application firewall, and content delivery network services.

7. Data-Center Management Issues (5 Marks)

1. User Satisfaction

Data centers must ensure **high-quality service** to a large user base and maintain operational effectiveness for **30+ years**.

2. Controlled Information Flow

Emphasis is on **streamlined communication, high availability (HA), and uninterrupted services.**

3. Multiuser Manageability

Must efficiently handle **traffic, database updates, and server maintenance** to support multiple users.

4. Scalability

The system should **scale storage, processing, and power** to handle growing workloads and data.

5. Reliability in Virtualization

Must integrate **failover, fault tolerance, and live VM migration** for disaster recovery and application continuity.

6. Security & Data Protection

Requires **strong security mechanisms** to protect against **network attacks**, ensuring **data privacy and integrity**.

8. Cloud Design Objectives (6 Marks)

1. Shift to Centralized Computing

Moves **processing, storage, and software** from personal desktops to **data centers** via the Internet.

2. Service Provisioning & Cloud Economics

Services are delivered based on **SLAs**, focusing on **efficiency** and **cost-effectiveness** (pay-as-you-go model).

3. Performance Scalability

Cloud systems must **scale efficiently** with growing **user demand** and workload.

4. Data Privacy Protection

Ensures **secure handling of sensitive data**, building **trust** in cloud storage and processing.

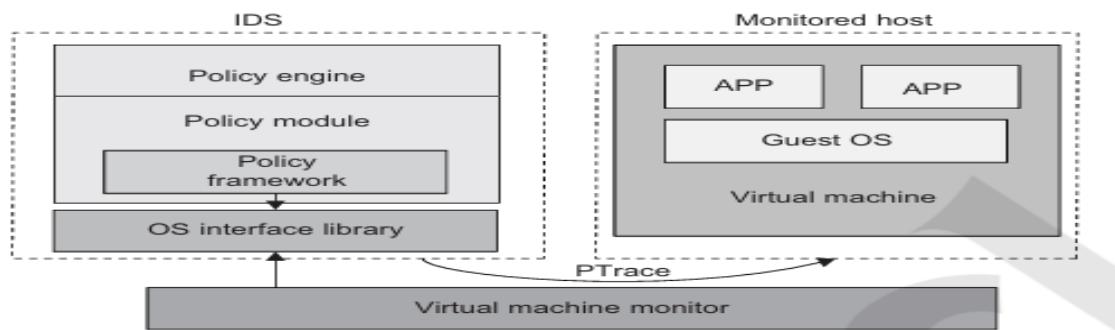
5. High Quality of Service (QoS)

Maintains **standardized performance** and **interoperability** across different providers.

6. Standards & Interfaces

Promotes **open APIs and protocols** to prevent **vendor lock-in** and enhance **portability**.

9. VM-Based Intrusion Detection



VM-Based Intrusion Detection (8 Marks)

1. Intrusion & IDS Overview

- Intrusion = **unauthorized access** by local or network users.
- **Intrusion Detection System (IDS)** monitors and detects such access based on behavioral patterns.
- IDS types:
 - **Host-based IDS (HIDS)**: Runs on individual systems, but can be attacked if the system is compromised.
 - **Network-based IDS (NIDS)**: Monitors traffic flow but may fail against **spoofed or fake actions**.

2. Role of Virtualization in IDS

- **VM-based IDS** isolates guest VMs, so an attack on one VM does **not affect others**—like NIDS.
- The **Virtual Machine Monitor (VMM)** can monitor access to **hardware and system software**, similar to HIDS.

3. Implementation Approaches

- IDS can run as:
 - An **independent process** inside each VM or in a **privileged VM**.
 - Or be **integrated into the VMM** with full access to hardware.

4. Policy Engine & Monitoring

- VM-based IDS includes a **policy engine and module**.
- Uses **OS interface libraries** and **PTrace** to trace and enforce security policies across guest VMs.

5. Intrusion Analysis & Logging

- Attack **logs** are used for post-intrusion analysis, but must be **tamper-proof**.
- IDS logging should run **independent of the compromised OS kernel** to maintain integrity.

6. Honeypots and Honeynets

- **Honeypots** are fake systems to **lure attackers** and observe behavior.
- Can be **physical or virtual**; in VMs, the host and VMM must remain secure to prevent real damage.
- Useful for both **protection** and **intrusion research**.

10. Advantages and Disadvantages of OS Extensions (6 Marks)

Advantages of OS Extensions

1. Low Overhead & High Scalability

- OS-level VMs have **minimal startup/shutdown time** and **low resource usage**, enabling **high scalability**.

2. Efficient Synchronization

- **VMs and host OS** can **synchronize state changes**, improving consistency and performance.

3. Shared Kernel

- All VMs **share a single OS kernel**, reducing duplication and improving resource utilization.

4. Flexible Access, Safe Isolation

- VMs can **access host resources** without modifying them, maintaining **security and isolation**.

Disadvantages of OS Extensions

1. Same OS Family Limitation

- All VMs must run **the same OS type** (e.g., Linux-based host can't run Windows VMs), reducing flexibility.

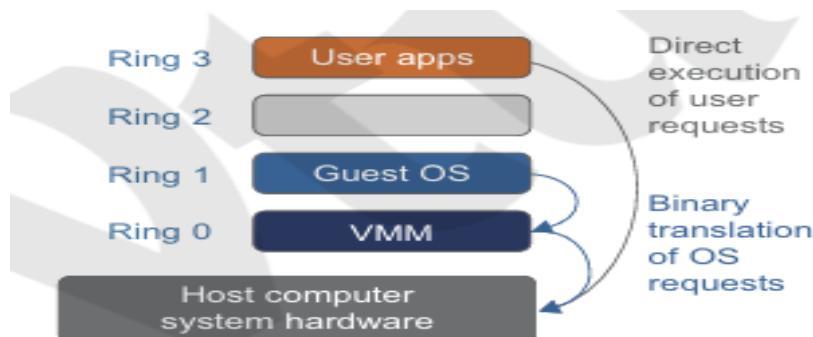
2. Limited OS Diversity

- In cloud environments with diverse user needs, **OS-level virtualization can't support multiple OS families simultaneously**.

3. Resource Duplication Overhead

- If each VM has **duplicated resources**, it leads to **increased overhead**, reducing the efficiency advantage.

11. Explain host based virtualisation



Host-Based Virtualization (Short Explanation - for 3 to 4 mark answers)

Host-based virtualization is a technique where the **virtualization layer (VM software)** is installed **on top of an existing host operating system**. The host OS manages the hardware, while **guest operating systems run inside VMs** created by the virtualization layer.

Key Features:

- **No host OS modification** is required to install virtualization.
 - The **host OS provides device drivers** and system services, simplifying VM design.
 - Both **guest OS applications** and **host OS applications** can run simultaneously.
-

Limitations:

- **Lower performance** compared to bare-metal hypervisors due to **four layers of mapping** during hardware access.
 - **Binary translation** is required when the guest OS uses a different instruction set, increasing overhead.
-

Use Case:

- Suitable for **desktop virtualization** or **testing multiple OSes** on personal systems, but **not ideal for high-performance cloud environments**.
-